

AI Model Routing Policy — Draft for Discussion

AI Dev Pipeline Forum · Digital Development · v0.1 · 8 July 2026 · Author: Stephen Cummins

TL;DR: One platform (Microsoft Foundry), model chosen per workload by data classification — GPT for personal/client data, Claude (ZDR) for internal non-personal work, Copilot Studio for assistants.

Principle

Microsoft Foundry is the single platform for programmatic AI model access: one Azure resource type, Entra ID authentication, Azure RBAC, consolidated billing (MACC-eligible). Model selection is a per-workload decision driven by data classification — not a one-vendor bet. This preserves optionality while giving governance a clear, auditable rule.

Routing tiers

Tier	Data classification	Route	Rationale
1	Personal data, client data, regulated or commercially sensitive content	Azure OpenAI (GPT), UK South / EU data zone	First-party Microsoft service: Microsoft operates inference, covered by existing Microsoft DPA and EU Data Boundary; UK residency available. No new processor.
2	Internal, non-personal data (templates, style rules, code, internal documentation)	Claude on Foundry — Hosted on Azure, Sweden Central, Zero Data Retention enabled	Best-fit model for agentic/coding/document workloads. Prompts and outputs processed on Azure infrastructure; ZDR means nothing retained post-call; Anthropic does not train on commercial data by default.
3	Conversational assistants (Q&A; over approved content)	Copilot Studio (existing route)	Already approved pattern (MaceyBot/ZoeBot); M365-native governance.

Compliance position (why Tier 2 is bounded)

- Claude in Foundry is a third-party Marketplace offering: Anthropic acts as independent data processor and sits outside Microsoft's standard DPA and EU Data Boundary. UK/EU data residency for Claude on Foundry is on the roadmap (expected 2026) but not yet available — hence the restriction to non-personal data.
- Mitigations in place for Tier 2: Hosted-on-Azure deployment (prompts/outputs processed on Azure infrastructure), Zero Data Retention, no training on commercial data by default, Anthropic DPA + SCCs, SOC 2 Type II / ISO 27001 / ISO 42001 certifications.
- Where a Tier 1 workload genuinely requires Claude-class capability, route via AWS Bedrock (EU) or Google Vertex AI (EU), which carry residency guarantees today — subject to separate approval.

Governance requirements (all tiers)

- Every AI workload registers in the AI Dev Pipeline tracker (intake → dev → test → launch → feedback) with data classification recorded at intake.
- DPIA + DPO sign-off required before any workload processes personal data; documented redaction/pseudonymisation where applicable.
- All model access via Entra ID and Azure RBAC; no standing API keys. Usage tracked through Azure cost management per workload.
- Deployment follows the established ALM pipeline (Dev → UAT → Prod) with change control.

Review triggers

- Claude UK/EU data zone reaches GA on Foundry → re-tier Tier 1 eligibility.
- Any change to Microsoft/Anthropic processor terms.
- Quarterly review at the AI Dev Pipeline forum.

Draft for discussion — subject to review by Mace Legal and the Data Protection Officer. This document summarises vendor documentation as at 8 July 2026 and is not legal advice.